

Introduction, Well-Posed Learning Problems, & Concept Learning

Relevant Readings: Sections 1.1, 2.1-2.2 in Mitchell

CS495 - Machine Learning, Fall 2009

Introductions

- ▶ Pass out syllabus

Introductions

- ▶ Pass out syllabus
- ▶ Why are we taking this course?

Introductions

- ▶ Pass out syllabus
- ▶ Why are we taking this course?

Machine Learning

- ▶ The study of *Machine learning* is about writing software that improves its own performance with experience

Machine Learning

- ▶ The study of *Machine learning* is about writing software that improves its own performance with experience
- ▶ Definition [Mitchell]: A computer program is said to learn from experience E with respect to some class of tasks T and performance measure P , if its performance at tasks in T , as measured by P , improves with experience E .

Machine Learning

- ▶ The study of *Machine learning* is about writing software that improves its own performance with experience
- ▶ Definition [Mitchell]: A computer program is said to learn from experience E with respect to some class of tasks T and performance measure P , if its performance at tasks in T , as measured by P , improves with experience E .

Driverless cars



Driverless cars



- ▶ Can phrase this as a well-posed learning problem

Driverless cars



- ▶ Can phrase this as a well-posed learning problem
 - ▶ T = driving a car on a busy interstate highway given image data from a forward-looking camera

Driverless cars



- ▶ Can phrase this as a well-posed learning problem
 - ▶ T = driving a car on a busy interstate highway given image data from a forward-looking camera
 - ▶ P = number of errors made, as judged by a human observer

Driverless cars



- ▶ Can phrase this as a well-posed learning problem
 - ▶ T = driving a car on a busy interstate highway given image data from a forward-looking camera
 - ▶ P = number of errors made, as judged by a human observer
 - ▶ E = image and control information recorded from driving sessions with a human driving the same vehicle

Driverless cars



- ▶ Can phrase this as a well-posed learning problem
 - ▶ T = driving a car on a busy interstate highway given image data from a forward-looking camera
 - ▶ P = number of errors made, as judged by a human observer
 - ▶ E = image and control information recorded from driving sessions with a human driving the same vehicle

What might you select for T , P , and E ?

- ▶ Handwriting recognition

What might you select for T , P , and E ?

- ▶ Handwriting recognition
- ▶ Chess playing program

What might you select for T , P , and E ?

- ▶ Handwriting recognition
- ▶ Chess playing program
- ▶ What are some other possible applications of machine learning?
 - ▶ How would you define the relevant learning problems involved?

What might you select for T , P , and E ?

- ▶ Handwriting recognition
- ▶ Chess playing program
- ▶ What are some other possible applications of machine learning?
 - ▶ How would you define the relevant learning problems involved?

Concept Learning

- ▶ Definition [Mitchell]: When the learning task is to infer a boolean-valued function from training examples of its input and output, this is known as *concept learning*.

Concept Learning

- ▶ Definition [Mitchell]: When the learning task is to infer a boolean-valued function from training examples of its input and output, this is known as *concept learning*.
 - ▶ So-named because you can think of it as predicting whether or not some property (or concept) holds

Concept Learning

- ▶ Definition [Mitchell]: When the learning task is to infer a boolean-valued function from training examples of its input and output, this is known as *concept learning*.
 - ▶ So-named because you can think of it as predicting whether or not some property (or concept) holds
- ▶ Although this may sound rather restrictive, concept learning is quite useful

Concept Learning

- ▶ Definition [Mitchell]: When the learning task is to infer a boolean-valued function from training examples of its input and output, this is known as *concept learning*.
 - ▶ So-named because you can think of it as predicting whether or not some property (or concept) holds
- ▶ Although this may sound rather restrictive, concept learning is quite useful

Concept Learning

- ▶ Examples of concept learning tasks:

Concept Learning

- ▶ Examples of concept learning tasks:
 - ▶ Movie suggestion
 - ▶ T = determine which movies a user will like
 - ▶ P = percent of unwatched movies correctly predicted (like/dislike)
 - ▶ E = examples of movies that the user likes and doesn't like

Concept Learning

- ▶ Examples of concept learning tasks:
 - ▶ Movie suggestion
 - ▶ T = determine which movies a user will like
 - ▶ P = percent of unwatched movies correctly predicted (like/dislike)
 - ▶ E = examples of movies that the user likes and doesn't like
 - ▶ Spam detection
 - ▶ T = determine which incoming emails are spam
 - ▶ $P = -10F_{pos} - F_{neg}$, where F_{pos} is the percentage of false positives and F_{neg} is the percentage of false negatives
 - ▶ E = examples of spam and nonspam emails

Concept Learning

- ▶ Examples of concept learning tasks:
 - ▶ Movie suggestion
 - ▶ T = determine which movies a user will like
 - ▶ P = percent of unwatched movies correctly predicted (like/dislike)
 - ▶ E = examples of movies that the user likes and doesn't like
 - ▶ Spam detection
 - ▶ T = determine which incoming emails are spam
 - ▶ $P = -10F_{pos} - F_{neg}$, where F_{pos} is the percentage of false positives and F_{neg} is the percentage of false negatives
 - ▶ E = examples of spam and nonspam emails
 - ▶ Graduate school admissions
 - ▶ T = predict which incoming grad students will eventually complete their PhD, based on their applications
 - ▶ P = the percentage of correct predictions
 - ▶ E = past grad students and their eventual grad school outcome

Concept Learning

- ▶ Examples of concept learning tasks:
 - ▶ Movie suggestion
 - ▶ T = determine which movies a user will like
 - ▶ P = percent of unwatched movies correctly predicted (like/dislike)
 - ▶ E = examples of movies that the user likes and doesn't like
 - ▶ Spam detection
 - ▶ T = determine which incoming emails are spam
 - ▶ $P = -10F_{pos} - F_{neg}$, where F_{pos} is the percentage of false positives and F_{neg} is the percentage of false negatives
 - ▶ E = examples of spam and nonspam emails
 - ▶ Graduate school admissions
 - ▶ T = predict which incoming grad students will eventually complete their PhD, based on their applications
 - ▶ P = the percentage of correct predictions
 - ▶ E = past grad students and their eventual grad school outcome